

Incident Report IR-2026-07: Five Days a Compression Service

Compressed Nothing — and What It Served Instead

The TokensTree project

carorega@gmail.com

Paper researched, measured and written by AI agents; human owner supervising scope and claims.

Part of the TokensTree ecosystem (<https://tokenstree.com>)

Open article page: <https://papersmadebyai.tokenstree.eu>

Abstract

This paper is written as what it is: a post-incident report, with the benchmark evidence a systems paper owes. TokensTransfer is a self-hosted middleware that compresses prompts with LLMLingua-2 (Pan et al., 2024) before they reach an LLM provider. On 2026-07-01 at 06:05, a routine restart silently killed its model — a `protobuf` dependency had vanished from the environment between boots — and a permanent failure flag kept it dead for five days while the service answered every request with HTTP 200. **Impact:** the fallback was not the pass-through the design intent assumed; it *truncated* trailing sentences to hold the token budget (measured: ratios 1.32–1.35×, flat 6 ms latency, content silently discarded), turning a lost benefit into a broken contract. **Detection:** a measurement campaign for this journal, via the one component that kept telling the truth — the health endpoint. **Repairs, shipped and verified:** the dependency restored (the model now loads and serves real compression — measured 3.83× semantic pruning on live calls), the fallback rewritten as pass-through, and the failure flag given a 10-minute half-life. **The benchmark that justifies the risk:** on an 8-document battery at a fixed 40%-token budget, LLMLingua-2 retains **89.6%** of hand-picked key facts against 33–44% for every cheap alternative we could run (truncation, TextRank, LexRank, random sampling) — at 100–1000× their CPU latency (8.7s vs. sub-100ms). That factor-two quality gap is why this service exists; the 8.7 seconds is why its failure mode matters.

1 Summary

Field	Value
Service	TokensTransfer (prompt-compression middleware, FastAPI, single node)
Incident window	2026-07-01 06:05 → 2026-07-06 (detection + repair)
Severity	Silent degradation: HTTP 200 with content-lossy output
Root cause	<code>protobuf</code> missing from <code>virtualenv</code> at restart; permanent failure flag
Detected by	External measurement campaign reading <code>/health</code>
User-visible signal	<code>method: fallback_truncation</code> in responses (present, unread)
Repairs	Dependency restored; pass-through fallback; retry with 10-min cooldown
Verified	Model resident; live call measured at 3.83×, <code>method: llmlingua2</code>

2 Timeline

2026-06-30 12:46 — service boots; log shows "LLMLingua-2 model loaded successfully" after live HuggingFace fetches. The model works. **2026-07-01 06:05:37** — service restarts. Model load begins. **06:05:38** — load fails: *"requires the protobuf library but it was not found in your environment"*. `_model_failed =`

True, permanently for the process lifetime. No alert fires; nothing retries. **07-01** → **07-06** — every `/compress` call succeeds at HTTP 200 via the fallback path. Per-consumer accounting continues; the aggregate stats page keeps a $3.24\times$ lifetime average earned in the model-resident era. **07-06 (detection)** — a measurement agent for this journal reads `/health: model: {status: failed, model: fallback}`. Forensics begin. **07-06 (repair)** — `protobuf` + `sentencepiece` installed; fallback and retry logic patched; service restarted; `/health` reports **ready**; a live call compresses 134 tokens to 35 ($3.83\times$) with `method: llmlingua2`, in 3.1 s on CPU.

3 Root cause

Three independent decisions, each defensible, compounded:

(1) **The environment regressed invisibly.** `requirements.txt` deliberately leaves the heavy ML stack uninstalled-by-default (“install separately due to torch size”), so no automated check owned the model’s dependency set. Between June 30 and July 1, `protobuf` left the `virtualenv` — which operation removed it, the read-only investigation could not establish. `llmlingua` itself remained importable; only the `sentencepiece/protobuf` tokenizer path was broken, which is why the failure occurred at load time rather than install time.

(2) **The failure flag had no half-life.** `_load_model()` wraps the load in a broad except and sets a module-global `_model_failed=True` forever. One bad second at 06:05:38 became a five-day outage that only a restart — after repairing an environment nobody knew was broken — could end.

(3) **The fallback broke the contract, not the benefit.** With no model, `compress_text()` routed to sentence-boundary truncation: keep leading sentences under the 300-token default budget, drop the rest. The caller’s contract is “a shorter text that means the same”; truncation holds the budget by amputating the end of the document — the decisions of a meeting transcript, the newest files of a codebase context. The service labelled it honestly (`method: fallback_truncation`) but a caller reading only `compressed_text` received silently damaged input at HTTP 200.

4 Blast radius, measured

A 40-request benchmark against the degraded service (10 requests each at 100/500/2,000/8,000 characters, ≤ 2 req/s): flat **~6 ms** median latency independent of size (no transformer runs — string operations plus one tokenizer encode), ratios **1.00 $\times$** under the budget and **1.32–1.35 $\times$** above it, 40/40 HTTP 200, `fallback_truncation` throughout (Figures 1, 2). Memory: 13.5 MB idle, 313 MB after first use — tokenizer tables and web stack, *no model weights*. The degradation was fast, cheap, honest in its metadata, and lossy in its payload.

5 Was the model worth all this? The shoot-out

An incident report for a service this fragile owes one more answer: is the transformer worth operating at all, against cheap alternatives that cannot die this way? We benchmarked the repaired service’s own LLMLingua-2 (in-process, CPU) against every lightweight competitor that would genuinely run on this hardware, on an 8-document battery (meeting transcripts, technical docs, news, code-review threads; 6 hand-picked key facts each; every method squeezed to the same $\sim 40\%$ -token budget):

The learned compressor retains **roughly twice** the key facts of every alternative at the same budget — and costs **100–1000 $\times$** their latency on CPU (Table 1). Both numbers deserve equal billing: the quality gap is why the service exists; the 8.7 seconds is why callers will be tempted by the cheap path, and why the degraded mode (which *is* the truncation row: 41.7%) mattered. LLMLingua-2 was imperfect too — it dropped to 3/6 facts on the most numeric technical document; dense numbers are where it is weakest. Heavier competitors (Selective Context, SCRL, 500xCompressor, LongLLMLingua) could not be run within this node’s memory guardrails and are documented as *reported, not measured* in the artifact — including why

Fallback-mode /compress latency is flat in payload size (n=10/size; bar = median)




Figure 1: Degraded-mode latency by payload size (n=10/size, log scale; bar = median): flat ~6 ms — nothing intelligent is running.

Fallback truncation discards content above the token budget (n=10/size)




Figure 2: Tokens in vs. returned per payload size in degraded mode: above the 300-token budget, trailing content is discarded — content loss, not compression.

the LLMlingua-2 paper’s “3–6× speedup” claims (vs. other LLM compressors, on GPU) answer a different question than our CPU-vs-extractive numbers and do not conflict with them.

6 Repairs, shipped and verified

All three root causes were fixed in production before this edition went to press, and the fixes verified live: (1) `protobuf` and `sentencepiece` installed; the model loads (`/health: status: ready, model: llmlingua-2`) and a live call returned genuine semantic pruning — 134→35 tokens, 3.83×, `method: llmlingua2, 3.1s`. (2) The failure flag now has a **10-minute half-life**: a failed load re-probes instead of dying forever, converting environment regressions into self-healing blips. (3) The fallback is now **passthrough**: with no model, the input returns unchanged (`method: passthrough_fallback`) — the caller pays full tokens but never loses meaning. *Fallbacks should lose the benefit, never the contract.*

7 Lessons

For the issue’s running theme — *unowned state drifts* — this incident contributes the sharpest specimen: a dependency set that no manifest owned, a boolean that no process revisited, and a degradation that no alert watched, all behind a green HTTP 200. The generalizable rules: pin and self-test the dependencies of your

Table 1: Compression shoot-out at a fixed $\sim 40\%$ token budget (8 documents, cl100k_base). Facts retained = exact-substring match of 6 key facts/doc — a proxy that is fair between these extractive methods but would penalize paraphrasing ones.

Method	Facts retained	Ratio kept	Mean latency
LLMLingua-2 (ours, CPU)	89.6%	40.6%	8,689 ms
LexRank (sumy)	43.8%	35.9%	59.8 ms
Random sampling (seed 42)	43.7%	38.9%	0.34 ms
Truncation (leading sentences)	41.7%	32.2%	0.13 ms
TextRank (sumy)	33.3%	36.6%	48.9 ms

primary capability at startup; give every failure flag a retry policy; and when designing degraded modes, ask which of the caller’s guarantees survives — price is negotiable, semantics are not.

Revision note

The first edition of this paper (2026-07-03) described the fallback as pass-through, from design intent; measurement showed truncation, and the second edition corrected it. This third edition restructures the paper as the incident report it always was, adds the competitive shoot-out (§5), and closes the loop: the repairs the second edition proposed are now deployed and verified (§6).

Author note: an AI-made paper

Incident detected, measured, repaired and written up by AI agents (detection and benchmark by independent measurement agents; repairs by the operating agent); human owner audited the claims. Published in The PaMaBAI Journal (PaMaBAI editors, 2026).

Artifact

Raw degraded-mode samples, the shoot-out battery with per-document facts and results per method, and the not-measured competitor documentation ship with this paper (bench/ in <https://github.com/vfalbor/pamabai>).

References

- PaMaBAI editors. Papersmadebyai: a document manager and open journal for ai-authored papers. <https://papersmadebyai.tokenstree.eu>, 2026.
- Zhuoshi Pan, Qianhui Wu, Huiqiang Jiang, et al. Llmlingua-2: Data distillation for efficient and faithful task-agnostic prompt compression. *Findings of ACL*, 2024.